

Context

- Jamboree has helped thousands of students like you make it to top colleges abroad. Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort.
- They recently launched a feature where students/learners can come to their website and check their probability of getting into the IVY league college. This feature estimates the chances of graduate admission from an Indian perspective.

Problem Statement :

- Your analysis will help Jamboree in understanding what factors are important in graduate admissions and how these factors are interrelated among themselves. It will also help predict one's chances of admission given the rest of the variables.

Column Profiling:

Serial No. (Unique row ID)
GRE Scores (out of 340)
TOEFL Scores (out of 120)
University Rating (out of 5)
Statement of Purpose and Letter of Recommendation Strength (out of 5)
Undergraduate GPA (out of 10)
Research Experience (either 0 or 1)
Chance of Admit (ranging from 0 to 1)

- Exploratory Data Analysis
- Linear Regression

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib import figure

import warnings
warnings.filterwarnings('ignore')

import statsmodels.api as sm
```

```
In [ ]:
```

```
In [2]: df = pd.read_csv("Jamboree_Admission.csv")
```

```
In [3]: df
```

Out[3]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65
...
495	496	332	108	5	4.5	4.0	9.02	1	0.87
496	497	337	117	5	5.0	5.0	9.87	1	0.96
497	498	330	120	5	4.5	5.0	9.56	1	0.93
498	499	312	103	4	4.0	5.0	8.43	0	0.73
499	500	327	113	4	4.5	4.5	9.04	0	0.84

500 rows x 9 columns

```
In [4]: data = df.copy()
```

```
In [ ]:
```

```
In [5]: # shape of the data
data.shape
```

```
Out[5]: (500, 9)
```

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score             500 non-null   int64
2   TOEFL Score           500 non-null   int64
3   University Rating     500 non-null   int64
4   SOP                   500 non-null   float64
5   LOR                   500 non-null   float64
6   CGPA                  500 non-null   float64
7   Research              500 non-null   int64
8   Chance of Admit       500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

```
In [7]: data.drop(["Serial No."],axis = 1, inplace = True)
```

```
In [8]: data.sample(5)
```

Out[8]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
109	304	103	5	5.0	4.0	8.64	0	0.68
185	327	113	4	4.5	4.5	9.11	1	0.89
264	325	110	2	3.0	2.5	8.76	1	0.75
413	317	101	3	3.0	2.0	7.94	1	0.49
33	340	114	5	4.0	4.0	9.60	1	0.90

```
In [ ]:
```

```
In [9]: # isnull ?
data.isna().sum()
```

Out[9]:

GRE Score	0
TOEFL Score	0
University Rating	0
SOP	0
LOR	0
CGPA	0
Research	0
Chance of Admit	0

dtype: int64

```
In [10]: # no null values found in data
```

```
In [ ]:
```

```
In [11]: data.columns
```

Out[11]:

```
Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA',
      'Research', 'Chance of Admit '],
      dtype='object')
```

```
In [12]: data.nunique()
```

Out[12]:

GRE Score	49
TOEFL Score	29
University Rating	5
SOP	9
LOR	9
CGPA	184
Research	2
Chance of Admit	61

dtype: int64

University Rating,SOP,LOR,Research are categorical variables.

all of the features are numeric , and ordinal . (University Rating,SOP,LOR,Research are discrete) and rest are continuous

```
In [ ]:
```

```
In [ ]:
```

Overall glance for correlations :

In [14]: `data.corr()`


Out[14]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
GRE Score	1.000000	0.827200	0.635376	0.613498	0.524679	0.825878	0.563398	0.810351
TOEFL Score	0.827200	1.000000	0.649799	0.644410	0.541563	0.810574	0.467012	0.792228
University Rating	0.635376	0.649799	1.000000	0.728024	0.608651	0.705254	0.427047	0.690132
SOP	0.613498	0.644410	0.728024	1.000000	0.663707	0.712154	0.408116	0.684137
LOR	0.524679	0.541563	0.608651	0.663707	1.000000	0.637469	0.372526	0.645365
CGPA	0.825878	0.810574	0.705254	0.712154	0.637469	1.000000	0.501311	0.882413
Research	0.563398	0.467012	0.427047	0.408116	0.372526	0.501311	1.000000	0.545871
Chance of Admit	0.810351	0.792228	0.690132	0.684137	0.645365	0.882413	0.545871	1.000000

In [15]: *# further correlation check is being done while Multicollinearity check for independent features and # correlation between independent and dependent features.*

In [16]: `plt.figure(figsize=(10,7))
sns.heatmap(data.corr(),annot = True,cmap = "Blues")`

Out[16]: `<AxesSubplot:>`



In []:

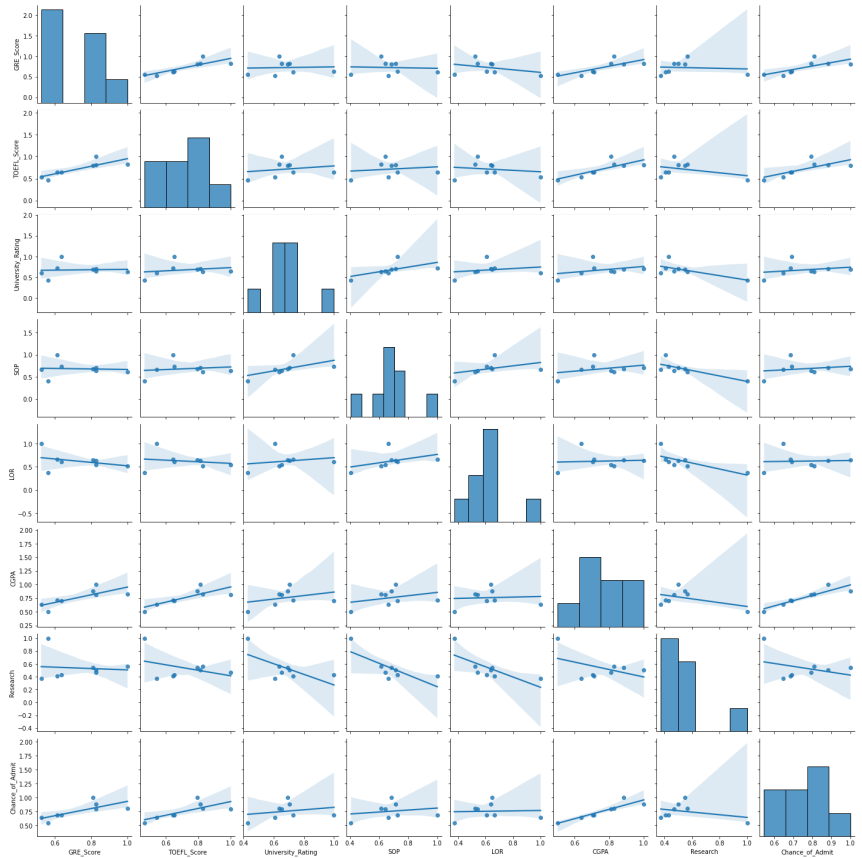
In [17]: `data.columns = ['GRE_Score', 'TOEFL_Score', 'University_Rating', 'SOP', 'LOR', 'CGPA',
'Research', 'Chance_of_Admit']`

In []:

pairplot , correlation and trend line with each variables:

```
In [18]: sns.pairplot(data.corr(),kind= 'reg',)
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x2ac1e717e80>
```



```
In [ ]:
```

check for outliers using IQR method

```
In [19]: def detect_outliers(data):
        length_before = len(data)
        Q1 = np.percentile(data,25)
        Q3 = np.percentile(data,75)
        IQR = Q3-Q1
        upperbound = Q3+1.5*IQR
        lowerbound = Q1-1.5*IQR
        if lowerbound < 0:
            lowerbound = 0

        length_after = len(data[(data>lowerbound)&(data<upperbound)])
        return f"{np.round((length_before-length_after)/length_before,4)} % Outliers data from input data found"
```

```
In [ ]:
```

```
In [ ]:
```

```
In [20]: for col in data.columns:
        print(col, " : ",detect_outliers(data[col]))
```

GRE_Score : 0.0 % Outliers data from input data found
TOEFL_Score : 0.0 % Outliers data from input data found
University_Rating : 0.0 % Outliers data from input data found
SOP : 0.0 % Outliers data from input data found
LOR : 0.024 % Outliers data from input data found
CGPA : 0.0 % Outliers data from input data found
Research : 0.44 % Outliers data from input data found
Chance_of_Admit : 0.004 % Outliers data from input data found

```
In [21]: # there are no significant amount of outliers found in the data
```

```
In [ ]:
```

```
In [ ]:
```


```
In [ ]:
```

Checking the distributions for Continuous Variables :

```
In [22]: # Chance_of_Admit
```


```
In [ ]:
```

```
In [23]: sns.distplot(data["Chance_of_Admit"])
sm.qqplot(data["Chance_of_Admit"],fit=True, line="45")
plt.show()
```



GRE_Score

```
In [24]: sns.distplot(data["GRE_Score"])
sm.qqplot(data["GRE_Score"],fit=True, line="45")
plt.show()
```




Chance of admit and GRE score are nearly normally distributed.

for EDA purpose , converting GRE score into bins , to check how distribution of chance of admit across the bins are :

```
In [28]: df["GRE_SCORE_CATEGORY"]=pd.qcut(df["GRE_Score"],20)
```

```
In [ ]:
```


```
In [29]: plt.figure(figsize=(14,5))
sns.boxplot(y = df["Chance of Admit "], x = df["GRE_SCORE_CATEGORY"])
plt.xticks(rotation = 90)
plt.show()
```



From above boxplot (distribution of chance of admission (probability of getting admission) as per GRE score) : with higher GRE score , there is high probability of getting an admission .

```
In [30]: sns.jointplot(data["GRE_Score"],data["Chance_of_Admit"], kind = "reg" )
```


```
Out[30]: <seaborn.axisgrid.JointGrid at 0x2ac23b33c70>
```




from above regression line| jointplot and boxlot we can observe a strong correlation of GRE score and chance of admit .

```
In [31]: # TOEFL_Score

sns.distplot(data["TOEFL_Score"])
sm.qqplot(data["TOEFL_Score"],fit=True, line="45")
plt.show()
plt.figure(figsize=(14,5))
sns.boxplot(y = data["Chance_of_Admit"], x = data["TOEFL_Score"])
```



```
Out[31]: <AxesSubplot:xlabel='TOEFL_Score', ylabel='Chance_of_Admit'>
```



Students having high toefl score, has higher probability of getting admission .


```
In [32]: data[["GRE_Score","TOEFL_Score","Chance_of_Admit"]].corr()
```

```
Out[32]:
```

	GRE_Score	TOEFL_Score	Chance_of_Admit
GRE_Score	1.000000	0.827200	0.810351
TOEFL_Score	0.827200	1.000000	0.792228
Chance_of_Admit	0.810351	0.792228	1.000000


```
In [33]: sns.jointplot(data["TOEFL_Score"],data["Chance_of_Admit"], kind = "reg" )
```



```
Out[33]: <seaborn.axisgrid.JointGrid at 0x2ac244bc190>
```



GRE_Score and Toefl_Score have very high correlation with Chance_of_Admit

```
In [ ]:
```

```
In [34]: # CGPA
sns.distplot(data["CGPA"])
sm.qqplot(data["CGPA"],fit=True, line="45")
plt.show()
```



```
In [35]: data[["CGPA", "Chance_of_Admit"]].corr()
```


Out[35]:

	CGPA	Chance_of_Admit
CGPA	1.000000	0.882413
Chance_of_Admit	0.882413	1.000000

CGPA also has a very high correlation with Chance of Admission

```
In [36]: sns.jointplot(data["CGPA"],data["Chance_of_Admit"], kind = "reg" )
```

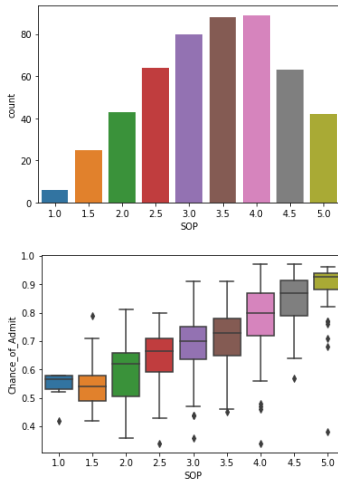
Out[36]: <seaborn.axisgrid.JointGrid at 0x2ac23aa2bb0>



GRE score, TOEFL score and CGPA has a strong correlation with chance of admission

```
In [ ]:
In [ ]:
In [ ]:
In [ ]: #CHECKING FOR REST OF THE FEATURES AND THEIR DISTRIBUTION :
In [ ]:
```

```
In [42]: # SOP strength
sns.countplot(data["SOP"])
plt.show()
sns.boxplot(y = data["Chance_of_Admit"], x = data["SOP"])
plt.show()
```

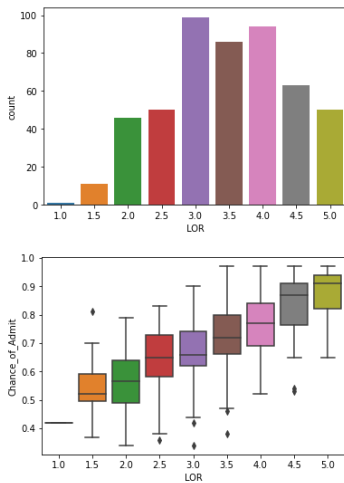


Distribution above shows , most occuring SOP strength us between 2.5 to 4.5

and having higher strength of SOP , bring more chance of getting admission !

```
In [ ]:
```

```
In [41]: sns.countplot(data["LOR"])
plt.show()
sns.boxplot(y = data["Chance_of_Admit"], x = data["LOR"])
plt.show()
```



Statement of Purpose and Letter of Recommendation Strength increases then the chances of admission also increases

```
In [43]: data[["SOP", "LOR", "Chance_of_Admit"]].corr()
```

```
Out[43]:
```

	SOP	LOR	Chance_of_Admit
SOP	1.000000	0.663707	0.684137
LOR	0.663707	1.000000	0.645365
Chance_of_Admit	0.684137	0.645365	1.000000

```
In [ ]:
```



```
In [ ]:
```

Distribution of Categorical variables

```
In [46]: data["University_Rating"].value_counts()
```

```
Out[46]: 3    162
         2    126
         4    105
         5     73
         1     34
         Name: University_Rating, dtype: int64
```



```
In [47]: sns.countplot(data["University_Rating"])
plt.show()
sns.boxplot(y = data["Chance_of_Admit"], x = data["University_Rating"])
plt.show()
```



higher the university rating , increase the chance of getting admission .

```
In [ ]: #Research
```

```
In [49]: sns.countplot(data["Research"])
plt.show()
sns.boxplot(y = data["Chance_of_Admit"], x = data["Research"])
plt.show()
```



for research student has higher chance of getting the admission.

In []:

In []:


Assumption check for Linear Regression :

In []:

```
In [50]: for col in data.columns[:-1]:
          print(col)
          plt.figure(figsize=(3,3))
          sns.jointplot(data[col],data["Chance_of_Admit"],kind="reg")
          plt.show()
```


GRE_Score

<Figure size 216x216 with 0 Axes>




TOEFL_Score

<Figure size 216x216 with 0 Axes>




University_Rating

<Figure size 216x216 with 0 Axes>




SOP

<Figure size 216x216 with 0 Axes>




LOR

<Figure size 216x216 with 0 Axes>




CGPA

<Figure size 216x216 with 0 Axes>



Research

<Figure size 216x216 with 0 Axes>



LOR, SOP , University rating and research are categorical variable, and amongst them chances of admits varies a lot.

```
In [ ]: # further assumption checks are done while building and testing model .
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

Regression using Sklearn library

Closed form solution technique for Linear Regression | OLS:

```
In [ ]:
```

```
In [ ]:
```

```
In [51]: X = data.drop(["Chance_of_Admit"],axis = 1)
y = data["Chance_of_Admit"]
```

```
In [ ]:
```

```
In [55]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
In [56]: model = LinearRegression()
```

```
In [57]: # train test splitting :
```

```
In [58]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [59]: model.fit(X_train,y_train)
```

```
Out[59]: LinearRegression()
```

```
In [61]: for idx, col in enumerate(X_train.columns):
print("Coefficient for {} is {}".format(col,model.coef_[idx]))
```

```
Coefficient for GRE_Score is 0.002134116998958902
Coefficient for TOEFL_Score is 0.0029507946431573742
Coefficient for University_Rating is 0.004842411688671617
Coefficient for SOP is 0.002095555922376041
Coefficient for LOR is 0.018600202256919177
Coefficient for CGPA is 0.11336157243184922
Coefficient for Research is 0.024713311522787978
```

```
In [62]: intercept = model.intercept_
intercept
```

```
Out[62]: -1.341760629850921
```

```
In [63]: # r2_score
model.score(X_test,y_test)
```

```
Out[63]: 0.7927524897595928
```

```
In [64]: # testing model on testing splitted data.
```

```
In [65]: y_pred = model.predict(X_test)
```

```
In [66]: print("MSE:", mean_squared_error(y_test, y_pred)) # MSE
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred))) # RMSE
print("MAE :", mean_absolute_error(y_test, y_pred) ) # MAE
print("r2_score:", r2_score(y_test, y_pred)) # r2score
```

```
MSE: 0.004429285498957574
RMSE: 0.0655287746564813
MAE : 0.04730057428620611
r2_score: 0.7927524897595928
```

since all the data is numeric and ordinal, keeping all the features , r₂ score is observed as 0.79 on test data .

```
In [ ]:
```

```
In [ ]:
```

Using Sklearn | Stochastic Gradient Descent Aalgorithm"

```
In [131]: X = data.drop(["Chance_of_Admit"],axis = 1)
y = data["Chance_of_Admit"]
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [132]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
In [133]: scaler.fit(X_train)
```

```
Out[133]: StandardScaler()
```

```
In [134]: X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test) # apply same transformation to test data
```

```
In [135]: from sklearn.linear_model import SGDRegressor

from sklearn.pipeline import make_pipeline
sgd = make_pipeline(StandardScaler(), SGDRegressor(max_iter=1000, tol=1e-3))
```

```
In [136]: sgd.fit(X_train, y_train)
```

```
Out[136]: Pipeline(steps=[('standardscaler', StandardScaler()),
                          ('sgdregressor', SGDRegressor())])
```

```
In [137]: y_pred = sgd.predict(X_test)
```

```
In [138]: y_test = y_test.values
```

```
In [139]: r2_score(y_test, y_pred)
```

```
Out[139]: 0.7903760694738095
```

```
In [ ]: # overserving very similar result as OLS .
# trying different algorithms and different variations with features.
```

```
In [ ]:
```

Linear Regression using Statsmodel library

```
In [164]: import statsmodels.api as sm
```

```
In [165]: X = data.drop(["Chance_of_Admit"],axis = 1)
y = data["Chance_of_Admit"]
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [166]: X_train_sm = X_train
X_test_sm = X_test
```

```
In [167]: X_train_sm = sm.add_constant(X_train_sm)
X_test_sm = sm.add_constant(X_test_sm)
```

```
In [168]: # added a constant in x_train , as stats model regression doesnt account for intercept separately
```

Multicollinearity check and further re-training model and testing :

```
In [169]: data.drop(["Chance_of_Admit"],axis = 1).corr()
```

```
Out[169]:
```


	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research
GRE_Score	1.000000	0.827200	0.635376	0.613498	0.524679	0.825878	0.563398
TOEFL_Score	0.827200	1.000000	0.649799	0.644410	0.541563	0.810574	0.467012
University_Rating	0.635376	0.649799	1.000000	0.728024	0.608651	0.705254	0.427047
SOP	0.613498	0.644410	0.728024	1.000000	0.663707	0.712154	0.408116
LOR	0.524679	0.541563	0.608651	0.663707	1.000000	0.637469	0.372526
CGPA	0.825878	0.810574	0.705254	0.712154	0.637469	1.000000	0.501311
Research	0.563398	0.467012	0.427047	0.408116	0.372526	0.501311	1.000000

```
In [ ]:
```

```
In [ ]:
```

```
In [170]: plt.figure(figsize=(10,7))
sns.heatmap(data.drop(["Chance_of_Admit"],axis = 1).corr(),annot = True,cmap = "Blues")
```

```
Out[170]: <AxesSubplot>
```



```
In [171]: # GRE score and Toefel score have a very high correlation with CGPA
# GRE score and TOEFL score also have a very high correlation
# CGPA and University Rating , SOP stength and CGPA, have a high correlation .
```

```
In [172]: # checking for Multicollinearity using vif score :
```

Variance Inflation Factor:

```
In [173]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [174]: vifs = []

for i in range(X_train_sm.shape[1]):

    vifs.append((variance_inflation_factor(exog = X_train_sm.values,
                                           exog_idx=i)))

pd.DataFrame({"coef_name": X_train_sm.columns,
              "vif": np.around(vifs,2)})
```

Out[174]:

	coef_name :	vif :
0	const	1571.81
1	GRE_Score	4.24
2	TOEFL_Score	4.06
3	University_Rating	2.59
4	SOP	2.71
5	LOR	1.98
6	CGPA	4.77
7	Research	1.47

```
In [175]: # VIF score are all below 5 , Look good , there doesnt seem significant multicollinearity.
```

```
In [176]: # model building
```

```
In [177]: olsres = sm.OLS(y_train,X_train_sm).fit()
```

```
In [178]: print(olsres.summary())
```

```

=====
                    OLS Regression Results
=====
Dep. Variable:      Chance_of_Admit    R-squared:                0.829
Model:              OLS                Adj. R-squared:          0.826
Method:             Least Squares      F-statistic:             272.1
Date:               Tue, 04 Oct 2022    Prob (F-statistic):      3.33e-146
Time:               10:20:32           Log-Likelihood:          573.41
No. Observations:   400                AIC:                   -1131.
Df Residuals:       392                BIC:                   -1099.
Df Model:           7
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.3418	0.116	-11.613	0.000	-1.569	-1.115
GRE_Score	0.0021	0.001	3.893	0.000	0.001	0.003
TOEFL_Score	0.0030	0.001	3.024	0.003	0.001	0.005
University_Rating	0.0048	0.004	1.185	0.237	-0.003	0.013
SOP	0.0021	0.005	0.428	0.669	-0.008	0.012
LOR	0.0186	0.005	4.131	0.000	0.010	0.027
CGPA	0.1134	0.011	10.633	0.000	0.092	0.134
Research	0.0247	0.007	3.476	0.001	0.011	0.039

```

=====
Omnibus:                 94.166    Durbin-Watson:              1.943
Prob(Omnibus):            0.000    Jarque-Bera (JB):           231.309
Skew:                    -1.158    Prob(JB):                   5.92e-51
Kurtosis:                 5.918    Cond. No.                   1.33e+04
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.33e+04. This might indicate that there are
strong multicollinearity or other numerical problems.
```

```
In [ ]:
```

```
In [180]: r2_score(y_test,olsres.predict(X_test_sm))
```

```
Out[180]: 0.7927524897595936
```

```
In [ ]: # same result of r2 value , as sklearn OLS regressor. ,
```

```
In [ ]:
```

Residual analysis :

```
In [181]: ypred = olsres.predict(X_train_sm)
```

```
In [182]: print("Mean of residuals : ",np.mean(y_train - ypred))
```


Mean of residuals : 1.1572687252936476e-15

```
In [183]: # distribution plot of all residuals
```


```
In [184]: Residuals = (y_train-ypred)
```

```
In [185]: sns.distplot(Residuals)
```

```
Out[185]: <AxesSubplot:ylabel='Density'>
```



```
In [186]: plt.scatter(y_train,Residuals)
plt.xlabel("Chances of Admit")
plt.ylabel("Residuals")
plt.axhline(y= 0)
plt.show()
```



Homoscedasticity

from above residual plot , we can observe the varinace is not so constant .

all residuals are not evenly distributed.

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```



```
In [198]: pd.DataFrame({ "coef_name": X_train_sm.columns,
                        "vif": np.around(vifs,2)})
```

```
Out[198]:
```

	coef_name	vif
0	const	1551.09
1	GRE_Score	3.68
2	TOEFL_Score	3.63
3	University_Rating	2.44
4	SOP	2.35
5	Research	1.45

compare to previous model , VIF score has improved

```
In [199]: olsres = sm.OLS(y_train,X_train_sm).fit()
print(olsres.summary())
```

```

=====
                    OLS Regression Results
=====
Dep. Variable:      Chance_of_Admit    R-squared:                0.761
Model:              OLS                Adj. R-squared:          0.758
Method:             Least Squares      F-statistic:              251.6
Date:               Tue, 04 Oct 2022    Prob (F-statistic):       3.30e-120
Time:              10:24:40             Log-Likelihood:           506.48
No. Observations:   400                 AIC:                    -1001.
Df Residuals:       394                 BIC:                    -977.0
Df Model:           5
Covariance Type:    nonrobust
=====
                    coef    std err          t      P>|t|    [0.025    0.975]
-----
const             -1.4911      0.135     -11.017    0.000     -1.757     -1.225
GRE_Score           0.0043      0.001      7.102    0.000      0.003     0.005
TOEFL_Score         0.0067      0.001      6.187    0.000      0.005     0.009
University_Rating   0.0168      0.005      3.597    0.000      0.008     0.026
SOP                 0.0206      0.005      3.830    0.000      0.010     0.031
Research            0.0326      0.008      3.901    0.000      0.016     0.049
=====
Omnibus:            77.416    Durbin-Watson:           1.864
Prob(Omnibus):      0.000    Jarque-Bera (JB):        138.987
Skew:               -1.094    Prob(JB):                 6.60e-31
Kurtosis:           4.884    Cond. No.                 1.32e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.32e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [ ]:
```

```
In [ ]:
```

```
In [200]: # re-training model with sklearn Library after dropping multicorrelated columns
```

```
In [201]: model = LinearRegression()
model.fit(X_train,y_train)
```

```
Out[201]: LinearRegression()
```

```
In [202]: for idx, col in enumerate(X_train.columns):
            print("Coefficient for {} is {}".format(col,model.coef_[idx]))
```

```

Coefficient for GRE_Score is 0.004275146280824592
Coefficient for TOEFL_Score is 0.006727403356408807
Coefficient for University_Rating is 0.01680793854723885
Coefficient for SOP is 0.02061850255251567
Coefficient for Research is 0.03256855858438903
```

```
In [203]: intercept = model.intercept_  
intercept
```

```
Out[203]: -1.4910580304577392
```

```
In [204]: model.score(X_test,y_test)
```

```
Out[204]: 0.7122332491254559
```

```
In [205]: mean_squared_error(y_test,y_pred) # MSE
```

```
Out[205]: 0.004480074258248521
```

```
In [206]: y_pred = model.predict(X_test)
```


```
In [ ]:
```

```
In [207]: r2_score(y_test,y_pred) # r2score
```

```
Out[207]: 0.7122332491254559
```


```
In [208]: sns.distplot((y_train.values-model.predict(X_train)))
```

```
Out[208]: <AxesSubplot:ylabel='Density'>
```



```
In [ ]:
```

```
In [210]: Residuals = (y_train - model.predict(X_train))  
plt.scatter(y_train,Residuals)  
plt.xlabel("Chances of Admit")  
plt.ylabel("Residuals")  
plt.axhline(y= 0)  
plt.show()
```



removing LOR and CGPA , retrained model gives R-2 values as 71% , which decreased compared to previous model .

```
In [ ]:
```

University rating, research are categorical data

trying one hot encoding on categorical data


```
In [214]: X = data.drop(["Chance_of_Admit"],axis = 1)
y = data["Chance_of_Admit"]
```

```
In [215]: X["University_Rating"] = X["University_Rating"].astype("str")
# X["SOP"] = X["SOP"].astype("str")
# X["LOR"] = X["LOR"].astype("str")
```

```
In [ ]:
```

```
In [216]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GRE_Score              500 non-null    int64
1   TOEFL_Score            500 non-null    int64
2   University_Rating      500 non-null    object
3   SOP                    500 non-null    float64
4   LOR                    500 non-null    float64
5   CGPA                   500 non-null    float64
6   Research               500 non-null    int64
dtypes: float64(3), int64(3), object(1)
memory usage: 27.5+ KB
```

```
In [217]: X = pd.get_dummies(X,columns=["University_Rating"], drop_first=True)
```

```
In [218]: X.sample(3)
```

```
Out[218]:
```

	GRE_Score	TOEFL_Score	SOP	LOR	CGPA	Research	University_Rating_2	University_Rating_3	University_Rating_4	University_Rating_5
181	305	107	2.5	2.5	8.42	0	1	0	0	0
36	299	106	4.0	4.0	8.40	0	1	0	0	0
484	317	106	3.5	3.0	7.89	1	0	1	0	0

```
In [219]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [220]: model = LinearRegression()
model.fit(X_train,y_train)
for idx, col in enumerate(X_train.columns):
    print("Coefficient for {} is {}".format(col,model.coef_[idx]))
```

```
Coefficient for GRE_Score is 0.0020879109282400535
Coefficient for TOEFL_Score is 0.0030751293420423244
Coefficient for SOP is 0.0022557230696422905
Coefficient for LOR is 0.018809911491126007
Coefficient for CGPA is 0.1129492493754535
Coefficient for Research is 0.024624133952513998
Coefficient for University_Rating_2 is -0.00743543800264374
Coefficient for University_Rating_3 is -0.008524187583490217
Coefficient for University_Rating_4 is -0.002728961462945642
Coefficient for University_Rating_5 is 0.014023745896441092
```

```
In [221]: intercept = model.intercept_
intercept
```

```
Out[221]: -1.3199903841650387
```

```
In [222]: y_predicted = model.predict(X_test)
r_2 = r2_score(y_test,y_predicted)
```

```
In [223]: print("MSE:",mean_squared_error(y_test,y_pred)) # MSE
print("RMSE:",np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
print("MAE:",mean_absolute_error(y_test,y_pred)) # MAE
print("r2_score : ",r2_score(y_test,y_predicted)) # r2score
```

```
MSE: 0.006150139489020719
RMSE: 0.07842282505126118
MAE: 0.058159525845528276
r2_score : 0.7925241207599244
```

```
In [224]: r_2
```

```
Out[224]: 0.7925241207599244
```

```
In [225]: model.score(X_test,y_test)
```

```
Out[225]: 0.7925241207599244
```

```
In [ ]:
```

```
In [226]: X_train_sm = X_train
X_train_sm = sm.add_constant(X_train_sm)
olsres = sm.OLS(y_train,X_train_sm).fit()

print(olsres.summary())
ypred = olsres.predict(X_train_sm)
Residuals = (y_train-ypred)
plt.scatter(y_train,Residuals)
plt.xlabel("Chances of Admit")
plt.ylabel("Residuals")
plt.axhline(y= 0)
plt.show()
```


OLS Regression Results						
=====						
Dep. Variable:	Chance_of_Admit		R-squared:	0.831		
Model:	OLS		Adj. R-squared:	0.827		
Method:	Least Squares		F-statistic:	191.2		
Date:	Tue, 04 Oct 2022		Prob (F-statistic):	2.13e-143		
Time:	10:27:08		Log-Likelihood:	575.33		
No. Observations:	400		AIC:	-1129.		
Df Residuals:	389		BIC:	-1085.		
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-1.3200	0.118	-11.186	0.000	-1.552	-1.088
GRE_Score	0.0021	0.001	3.793	0.000	0.001	0.003
TOEFL_Score	0.0031	0.001	3.142	0.002	0.001	0.005
SOP	0.0023	0.005	0.460	0.646	-0.007	0.012
LOR	0.0188	0.005	4.178	0.000	0.010	0.028
CGPA	0.1129	0.011	10.577	0.000	0.092	0.134
Research	0.0246	0.007	3.452	0.001	0.011	0.039
University_Rating_2	-0.0074	0.013	-0.554	0.580	-0.034	0.019
University_Rating_3	-0.0085	0.014	-0.599	0.550	-0.037	0.019
University_Rating_4	-0.0027	0.017	-0.163	0.871	-0.036	0.030
University_Rating_5	0.0140	0.019	0.752	0.452	-0.023	0.051
=====						
Omnibus:	89.185	Durbin-Watson:	1.944			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	210.672			
Skew:	-1.113	Prob(JB):	1.79e-46			
Kurtosis:	5.772	Cond. No.	1.36e+04			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.36e+04. This might indicate that there are strong multicollinearity or other numerical problems.



```
In [227]: vifs = []

for i in range(X_train_sm.shape[1]):
    vifs.append((variance_inflation_factor(exog = X_train_sm.values,
                                           exog_idx=i)))

pd.DataFrame({"coef_name": X_train_sm.columns,
              "vif": np.around(vifs,2)})
```

Out[227]:

	coef_name :	vif :
0	const	1642.70
1	GRE_Score	4.29
2	TOEFL_Score	4.10
3	SOP	2.71
4	LOR	1.98
5	CGPA	4.79
6	Research	1.48
7	University_Rating_2	4.16
8	University_Rating_3	5.23
9	University_Rating_4	5.15
10	University_Rating_5	5.37

```
In [ ]: # Converting University rating into category |and applying one hot encoding ,
# Multicollinearity seems to be increasing.
# though r_2 value is increased.
```

```
In [ ]:
```

```
In [ ]:
```

```
In [228]: # retraining after RFE :
```

recursive feature elimination (RFE) to select features :

```
In [229]: data
```

Out[229]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Chance_of_Admit
0	337	118	4	4.5	4.5	9.65	1	0.92
1	324	107	4	4.0	4.5	8.87	1	0.76
2	316	104	3	3.0	3.5	8.00	1	0.72
3	322	110	3	3.5	2.5	8.67	1	0.80
4	314	103	2	2.0	3.0	8.21	0	0.65
...
495	332	108	5	4.5	4.0	9.02	1	0.87
496	337	117	5	5.0	5.0	9.87	1	0.96
497	330	120	5	4.5	5.0	9.56	1	0.93
498	312	103	4	4.0	5.0	8.43	0	0.73
499	327	113	4	4.5	4.5	9.04	0	0.84

500 rows × 8 columns

```
In [241]: from sklearn.feature_selection import RFE
```

```
In [242]: LRm = LinearRegression()
```

```
In [243]: rfe = RFE(LRm,n_features_to_select=5)
```

```
In [244]: rfe
```

```
Out[244]: RFE(estimator=LinearRegression(), n_features_to_select=5)
```

```
In [245]: X = data.drop(["Chance_of_Admit"],axis = 1)
y = data["Chance_of_Admit"]
```

```
In [246]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
In [247]: rfe = rfe.fit(X_train,y_train)
```

```
In [248]: rfe.support_
```

```
Out[248]: array([False,  True,  True, False,  True,  True,  True])
```

```
In [249]: data.columns
```

```
Out[249]: Index(['GRE_Score', 'TOEFL_Score', 'University_Rating', 'SOP', 'LOR', 'CGPA',
               'Research', 'Chance_of_Admit'],
              dtype='object')
```

```
In [250]: rfe.ranking_
```

```
Out[250]: array([2, 1, 1, 3, 1, 1, 1])
```

```
In [251]: X_train.columns
```

```
Out[251]: Index(['GRE_Score', 'TOEFL_Score', 'University_Rating', 'SOP', 'LOR', 'CGPA',
               'Research'],
              dtype='object')
```

```
In [252]: X = data.drop(["Chance_of_Admit","University_Rating"],axis = 1)
y = data["Chance_of_Admit"]
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)

X_train_sm = X_train
X_test_sm = X_test

X_train_sm = sm.add_constant(X_train_sm)
X_test_sm = sm.add_constant(X_test_sm)
```

```
In [253]: X_train_sm = X_train
X_train_sm = sm.add_constant(X_train_sm)
olsres = sm.OLS(y_train,X_train_sm).fit()

print(olsres.summary())
ypred = olsres.predict(X_train_sm)
Residuals = (y_train-ypred)
plt.scatter(y_train,Residuals)
plt.xlabel("Chances of Admit")
plt.ylabel("Residuals")
plt.axhline(y= 0)
plt.show()
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Chance_of_Admit    R-squared:                0.829
Model:                  OLS                Adj. R-squared:           0.826
Method:                 Least Squares       F-statistic:              316.9
Date:                   Tue, 04 Oct 2022     Prob (F-statistic):       3.61e-147
Time:                   10:31:08             Log-Likelihood:           572.70
No. Observations:       400                 AIC:                     -1131.
Df Residuals:           393                 BIC:                     -1103.
Df Model:                6
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.3759	0.112	-12.291	0.000	-1.596	-1.156
GRE_Score	0.0022	0.001	3.985	0.000	0.001	0.003
TOEFL_Score	0.0030	0.001	3.099	0.002	0.001	0.005
SOP	0.0042	0.005	0.906	0.366	-0.005	0.013
LOR	0.0194	0.004	4.353	0.000	0.011	0.028
CGPA	0.1154	0.011	10.960	0.000	0.095	0.136
Research	0.0252	0.007	3.544	0.000	0.011	0.039

```


=====
Omnibus:                 94.122    Durbin-Watson:           1.948
Prob(Omnibus):            0.000    Jarque-Bera (JB):         231.319
Skew:                     -1.157    Prob(JB):                 5.88e-519
Kurtosis:                  5.920    Cond. No.                  1.28e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.28e+04. This might indicate that there are strong multicollinearity or other numerical problems.



```
In [254]: r2_score(y_test,olsres.predict(X_test_sm))
```

```
Out[254]: 0.7909249818462333
```

```
In [ ]:
```

```
In [ ]:
```

Inferences and Recommendations :

```
In [ ]:
```

Basic EDA and structure of data :

- First column was observed as unique row identifier which was dropped and wasn't required for neither EDA or modeling in our case.
- University Rating,SOP,LOR,Research are categorical variables. (still ordinal , have used as it is for model training.)
- all of the features are numeric , and ordinal . (University Rating,SOP,LOR,Research are discrete) and rest are continuous
- further correlation check is being done while Multicollinearity check for independent features and correlation between independent and dependent features.
- There were no significant amount of outliers found in the data.

